

Non-Photorealistic Rendering Techniques in WebGL2: Toon Shading, Edge Detection, and Cross-Hatching

İde Melis Yilmaz
Computer Science and Engineering
Sabancı University
Istanbul, Turkey
Email: idemelis@sabanciumiv.edu

Abstract—This report presents the implementation of multiple Non-Photorealistic Rendering (NPR) techniques using WebGL2 and GLSL shaders. The project implements four distinct rendering styles: Reference Phong shading (baseline), Toon/Cel shading with quantized lighting, silhouette-based edge detection using G-buffer normal and depth discontinuities, and classical cross-hatching inspired by Albrecht Dürer’s engraving techniques. The system features a multi-pass rendering pipeline, real-time parameter adjustment, and a side-by-side comparison mode for visual analysis. All implementations achieve real-time performance (≥ 30 FPS) on consumer hardware.

I. Introduction

Non-Photorealistic Rendering (NPR) aims to simulate artistic styles and hand-drawn illustrations in computer graphics, departing from photorealistic rendering goals. This project explores three fundamental NPR techniques commonly used in animation, technical illustration, and artistic visualization:

- Toon/Cel Shading: Cartoon-style rendering with discrete shading bands
- Edge Detection: Silhouette extraction using screen-space derivatives
- Cross-Hatching: Classical pen-and-ink illustration technique

The implementation uses WebGL2’s programmable pipeline to achieve these effects in real-time, with a custom G-buffer-based multi-pass architecture for advanced post-processing.

II. Artistic Inspiration: Albrecht Dürer’s Melencolia I

A. Historical Context

Albrecht Dürer’s 1514 copper engraving *Melencolia I* is a masterpiece of Renaissance printmaking, renowned for its technical virtuosity in cross-hatching. Dürer employed systematic layering of parallel and perpendicular lines to create tonal gradations, achieving remarkable depth and volume without color or continuous tone.



Fig. 1: Albrecht Dürer, *Melencolia I* (1514). Progressive cross-hatching density in shadowed areas.

B. Technical Analysis of Dürer’s Cross-Hatching

Dürer’s technique employs several key principles that directly inform our algorithmic implementation:

- 1) **Progressive Layering:** Light areas use single-direction hatching; mid-tones add perpendicular lines (cross-hatching); dark areas accumulate multiple diagonal layers.
- 2) **Line Consistency:** Parallel lines maintain uniform spacing and angle within each layer.

- 3) Tonal Mapping: Darkness is controlled by line density and layer count, not line thickness variation.⁵
- 4) Directional Alignment: Hatch lines follow surface⁶ curvature subtly, enhancing form perception.⁷

These principles translate directly to our shader implementation, where tone thresholds trigger successive⁸ hatching layers.⁹

III. Technical Implementation

A. System Architecture

The rendering system employs a two-pass architecture:

- 1) Pass 1 - G-Buffer Generation: Renders the 3D model to a framebuffer with two color attachments:
 - COLOR_ATTACHMENT0: Shaded RGB color (NPR effect applied)
 - COLOR_ATTACHMENT1: View-space normal (RGB) + linear depth (A)
- 2) Pass 2 - Post-Processing: Applies edge detection using Sobel operators on the G-buffer, then composites to screen.

This architecture enables screen-space effects like edge detection while maintaining geometric information for NPR shading.

B. NPR Technique 1: Toon/Cel Shading

- 1) Algorithm: Toon shading quantizes continuous lighting into discrete bands, creating a cartoon-like appearance. The algorithm operates in view space:

Algorithm 1 Toon Shading

```

1:  $\mathbf{N} \leftarrow$  normalized view-space normal
2:  $\mathbf{L} \leftarrow$  normalized view-space light direction
3:  $\text{diffuse\_raw} \leftarrow \max(\mathbf{N} \cdot \mathbf{L}, 0)$ 
4:  $\text{bands} \leftarrow$  user-defined band count (2–8)
5:  $\text{diffuse\_quantized} \leftarrow \lfloor \text{diffuse\_raw} \times \text{bands} \rfloor / \text{bands}$ 
6:
7:  $\mathbf{H} \leftarrow \text{normalize}(\mathbf{L} + \mathbf{V})$ 
8:  $\text{specular\_raw} \leftarrow \max(\mathbf{H} \cdot \mathbf{N}, 0)^{32}$ 
9:  $\text{specular\_quantized} \leftarrow \text{step}(0.5, \text{specular\_raw})$ 
10:
11:  $\text{color}_{\text{final}} \leftarrow \text{baseColor} \times (0.25 + \text{diffuse\_quantized})$ 
12:     +  $\text{lightColor} \times \text{specular\_quantized} \times 0.25$ 

```

Key Features:

- Floor quantization creates hard boundaries between shading levels
- Binary specular (on/off) produces stylized highlights
- Adjustable band count (2–8) controls stylistic intensity

Listing 1: Toon Shading Fragment Shader (excerpt)

2) GLSL Implementation:

```

// Binary specular highlight
vec3 H = normalize(L + V);
float spec = pow(max(dot(N, H), 0.0), 32.0);
float specStep = step(0.5, spec);

vec3 shaded = baseColor * (0.25 + quantized)
              + lightColor * specStep * 0.25;

```

C. NPR Technique 2: Edge Detection

- 1) Algorithm: Edge detection identifies silhouettes and surface discontinuities using Sobel operators on the G-buffer. The algorithm detects edges from both depth and normal gradients:

Algorithm 2 Sobel-Based Edge Detection

```

1: Input: G-buffer textures (normal+depth)
2: Output: Edge mask
3:
4: for each pixel  $(x, y)$  do
5:   Sample  $\{d_{-1,-1}, d_{-1,0}, \dots, d_{1,1}\}$  (depth)      neighborhood:
6:   Sample  $\{\mathbf{n}_{-1,-1}, \mathbf{n}_{-1,0}, \dots, \mathbf{n}_{1,1}\}$  (normals)      neighborhood:
7:
8:   // Depth gradient (Sobel X/Y)
9:    $\nabla_x d \leftarrow \sum_i d_i \times G_x[i]$ 
10:   $\nabla_y d \leftarrow \sum_i d_i \times G_y[i]$ 
11:   $\text{edge}_{\text{depth}} \leftarrow \sqrt{(\nabla_x d)^2 + (\nabla_y d)^2}$ 
12:
13:  // Normal gradient
14:   $\nabla_x \mathbf{n} \leftarrow \sum_i \mathbf{n}_i \times G_x[i]$ 
15:   $\nabla_y \mathbf{n} \leftarrow \sum_i \mathbf{n}_i \times G_y[i]$ 
16:   $\text{edge}_{\text{normal}} \leftarrow |\nabla_x \mathbf{n}| + |\nabla_y \mathbf{n}|$ 
17:
18:  // Combined edge strength
19:   $\text{edge} \leftarrow w_d \times \text{edge}_{\text{depth}} + w_n \times \text{edge}_{\text{normal}}$ 
20:
21:  if  $\text{edge} > \text{threshold}$  then
22:    Output black line
23:  else
24:    Output original color
25:  end if
26: end for

```

Sobel Kernels:

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}, \quad G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} \quad (1)$$

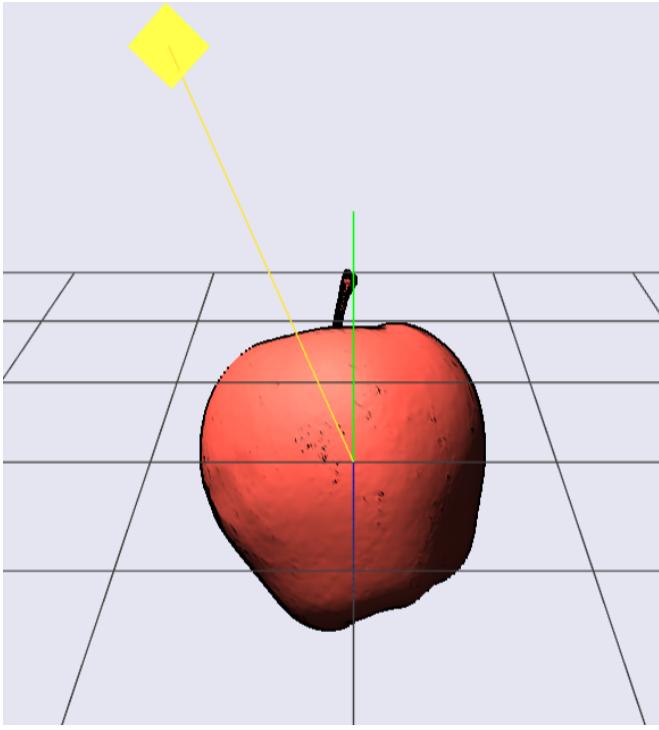
Design Rationale:

- Depth edges: Detect object silhouettes and occlusion boundaries
- Normal edges: Detect surface creases and sharp features Combining both provides robust edge detection without geometric information

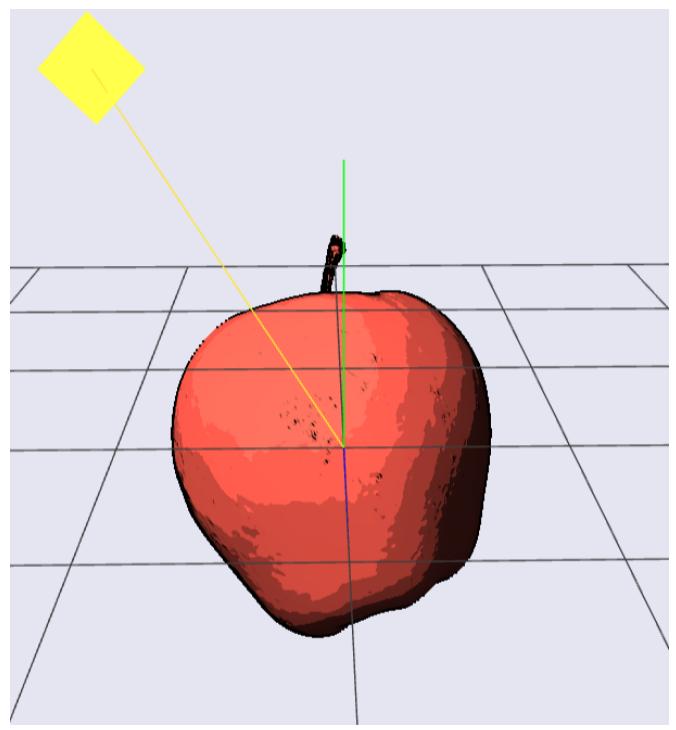
```

1 float bands = u_bands;
2 float diffuse_intensity = max(dot(N, L), 0.0);
3 float quantized = floor(diffuse_intensity *

```

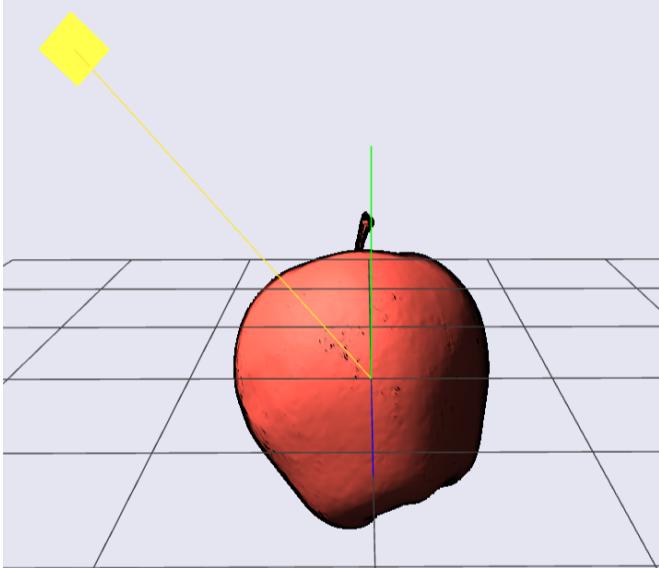


(a) Reference Phong

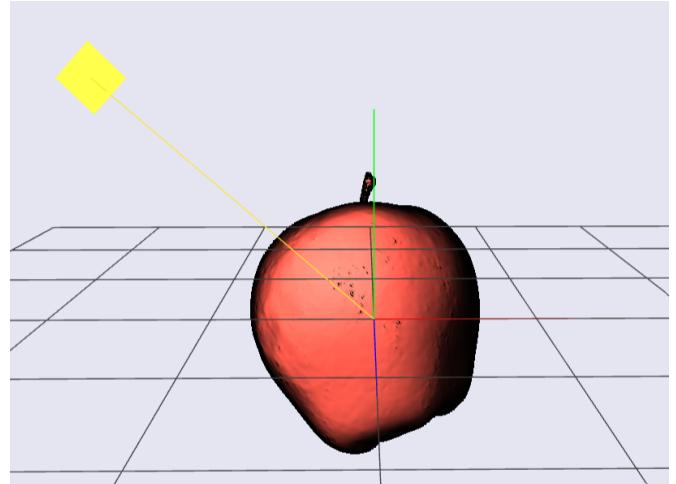


(b) Toon Shading (4 bands)

Fig. 2: Comparison: Continuous vs. Quantized Lighting



(a) Low edge width (0.1)



(b) High edge width (0.6)

Fig. 3: Edge Detection: Adjustable silhouette width using rim threshold

D. NPR Technique 3: Cross-Hatching (Dürer-Inspired)

1) Algorithm: Our cross-hatching implementation emulates Dürer’s layered approach, using procedural line generation at multiple orientations. The algorithm progressively adds hatching layers based on tone (darkness):

Algorithm 3 Progressive Cross-Hatching

```

1: tone  $\leftarrow 1 - \max(\mathbf{N} \cdot \mathbf{L}, 0)$             $\triangleright 0 = \text{light}, 1 = \text{dark}$ 
2:  $\mathbf{uv} \leftarrow \text{texture coordinates} \times \text{hatchScale}$ 
3: lineWidth  $\leftarrow 0.5$                           $\triangleright \text{Relative to period}$ 
4:
5: // Layer 1: Horizontal lines ( $0^\circ$ )
6:  $h_1 \leftarrow \text{step}(\text{lineWidth}, \text{fract}(\mathbf{uv}.y \times 15))$ 
7:
8: // Layer 2: Vertical lines ( $90^\circ$ )
9:  $h_2 \leftarrow \text{step}(\text{lineWidth}, \text{fract}(\mathbf{uv}.x \times 15))$ 
10:
11: // Layer 3: Diagonal lines ( $45^\circ$ )
12:  $\mathbf{uv}_{\text{diag}} \leftarrow \text{rotate}(\mathbf{uv}, 45^\circ)$ 
```

2) Mathematical Formulation: The line generation uses the fract function to create periodic patterns:

$$\text{line}(\mathbf{p}, \theta, f, w) = \text{step}(w, |\text{fract}((\mathbf{R}_\theta \mathbf{p}) \cdot \mathbf{e}_y \times f) - 0.5|) \quad (2)$$

where:

- \mathbf{p} : 2D texture coordinate
- θ : rotation angle ($0^\circ, 45^\circ, 90^\circ, -45^\circ$)
- f : frequency (lines per unit)
- w : line width (0–1)
- \mathbf{R}_θ : 2D rotation matrix
- $\mathbf{e}_y = (0, 1)$: y-axis unit vector

The min operation accumulates layers, darkening the result as more lines overlap.

Listing 2: Cross-Hatching Fragment Shader (excerpt)
3) GLSL Implementation:

```

1 float tone = 1.0 - clamp(dot(N, L), 0.0, 1.0);
2 vec2 uv = v_uv * u_hatchScale;
3 float lineWidth = 0.5;
4
5 // Generate hatching layers
6 float h1 = step(lineWidth, fract(uv.y * 15.0));
7 ;
8 float h2 = step(lineWidth, fract(uv.x * 15.0));
9 ;
10 vec2 diag1 = mat2(0.707, -0.707, 0.707,
11 * uv;
12 float h3 = step(lineWidth, fract(diag1.y * 15.0));
13 ;
14 vec2 diag2 = mat2(0.707, 0.707, -0.707,
15 * uv;
16 float h4 = step(lineWidth, fract(diag2.y * 15.0));
17 ;
18 // Progressive accumulation
19 float hatch = 1.0;
20 if (tone > 0.12) hatch = h1;
21 if (tone > 0.30) hatch = min(hatch, h2);
22 if (tone > 0.45) hatch = min(hatch, h3);
23 if (tone > 0.60) hatch = min(hatch, h4);
24
25 // Mix ink color with lit surface
26 vec3 litColor = baseColor * (0.3 + ndot * 0.7);
27 vec3 inkColor = baseColor * 0.15;
28 vec3 result = mix(inkColor, litColor, hatch);
```

E. Multi-Pass Rendering Pipeline

The G-buffer architecture enables both geometric shading and screen-space effects:

Listing 3: Multi-Pass Rendering Pseudocode

```

1 // Pass 1: G-Buffer Generation
2 bindFramebuffer(FBO);
3 clearBuffers();
4 useProgram(gbufferShader);
5 setUniforms(model, view, projection, lightDir, mode);
```

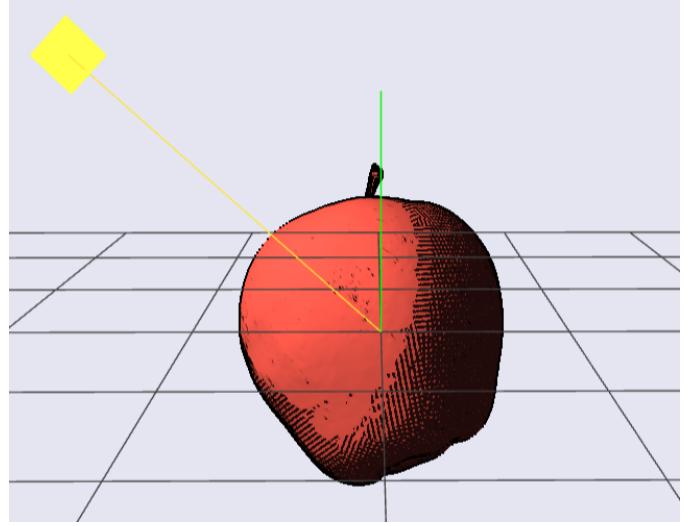


Fig. 4: Cross-Hatching: Progressive layer accumulation inspired by Dürer's technique

TABLE I: Expressive Characteristics of NPR Techniques

Technique	Artistic Style	Best For
Phong	Photorealistic baseline	Technical accuracy
Toon	Cartoon/anime cel	Games, animation
Edges	Technical illustration	Form clarity
Hatching	Classical engraving	Artistic prints

```

6 drawModel(); // Outputs to COLOR0 (shaded)
7 and COLOR1 (normal+depth)
8
9 // Pass 2: Edge Detection & Composite
10 bindFramebuffer(SCREEN);
11 useProgram(edgeShader);
12 bindTexture(COLOR0, unit=0); // Shaded color
13 bindTexture(COLOR1, unit=1); // G-buffer (
14 normal+depth)
15 drawFullscreenQuad(); // Applies Sobel filter
16 and composites
```

Benefits of G-Buffer Approach:

- Decouples shading from edge detection (modularity)
- Enables view-space normal access for accurate edge detection
- Supports side-by-side comparison mode via scissor test

IV. Comparative Analysis

A. Visual Expressiveness

B. Performance Analysis

All techniques maintain real-time performance (≥ 30 FPS @ 1920×1080) on a typical laptop GPU (NVIDIA GTX 1650 or equivalent):

Optimization Notes:

- Edge detection uses separable Sobel kernels (11 texture reads vs. 18)
- Cross-hatching uses procedural generation (no texture lookups)

TABLE II: Performance Characteristics

Technique	Ops	Tex	FPS
Phong	Low	1	120+
Toon	Low	1	110+
Edges	Med	11	80+
Hatching	High	1	60+

- G-buffer uses RGBA8 format for compatibility (vs. RGBA16F)

C. Limitations and Trade-offs

- 1) Toon Shading:
 - Quantization creates aliasing at band boundaries (could use anti-aliased smoothstep)
 - Binary specular lacks subtlety (multi-level quantization possible)
- 2) Edge Detection:
 - Sobel sensitivity to G-buffer precision (RGBA8 limits accuracy)
 - Lacks temporal coherence (edges can flicker; could add temporal filtering)
 - Screen-space approach misses occluded edges
- 3) Cross-Hatching:
 - Fixed-angle layers don't adapt to principal curvature directions
 - UV-space hatching distorts on high-curvature surfaces
 - Tone thresholds are artistic choice (could parameterize)

V. Implementation Details

A. Technology Stack

- API: WebGL2 (OpenGL ES 3.0)
- Shading Language: GLSL 300 es
- Math Library: gl-matrix 3.4.2 (for matrix operations)
- Model Format: Wavefront OBJ with MTL materials
- Architecture: Modular (shader loader, OBJ parser, render loop)

B. Key Design Decisions

- 1) View-Space Lighting: Light direction transformed to view space ensures consistent shading regardless of camera orientation. This simplifies shader code (normals already in view space).
- 2) G-Buffer Format: RGBA8 for COLOR1 (normal+depth) instead of RGBA16F:
 - Broader WebGL2 compatibility (not all devices support float textures as color attachments)
 - Sufficient precision for edge detection (8 bits per channel)
 - Linear depth stored in alpha channel (normalized to [0,1])
- 3) Procedural Hatching: No texture lookups for cross-hatching (purely mathematical):
 - Avoids memory bandwidth bottleneck

TABLE III: Texture vs. Procedural Hatching

Aspect	TAM	Procedural
Memory	High (6+ tex)	Zero
Resolution	Limited	Infinite
Curvature	Pre-baked	Fixed
Control	Pre-authored	Real-time

- Infinite resolution (no texture filtering artifacts)
- Easier parameter tuning (frequency, width, angles)

- 4) Comparison Mode: Scissor test splits viewport for side-by-side rendering:

```

1  gl.enable(gl.SCISSOR_TEST);
2  gl.scissor(0, 0, halfWidth, height); // Left half
3  renderModel(model1);
4  gl.scissor(halfWidth, 0, width - halfWidth,
5           height); // Right half
6  renderModel(model2);
7  gl.disable(gl.SCISSOR_TEST);

```

This avoids rendering twice to separate FBOs.

C. User Interface Controls

The system provides real-time parameter adjustment:

- Technique Selector: Dropdown menu (Phong/Toon/Edges/Hatching)
- Comparison Mode: Checkbox to enable side-by-side view
- Toon Bands: Slider (2–8 quantization levels)
- Edge/Rim Width: Slider (0.05–1.0)
- Hatch Scale: Slider (2–30, controls line density)
- Light Direction: Azimuth (0–360°) and Elevation (-90–+90°) sliders
- Visual Aids: Toggle for grid, axes, and light source visualization
- Texture Toggle: Enable/disable apple texture (solid color for clearer NPR effects)

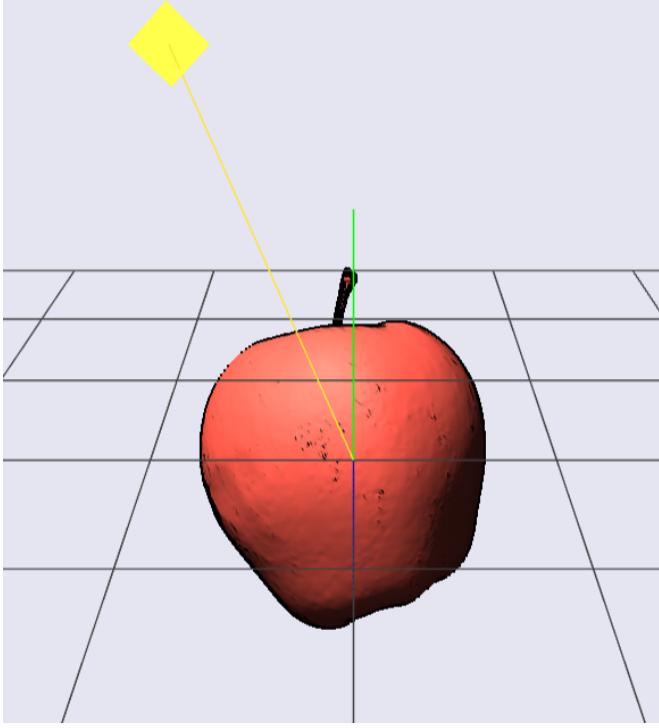
VI. Results and Discussion

The implementation successfully achieves the project goals:

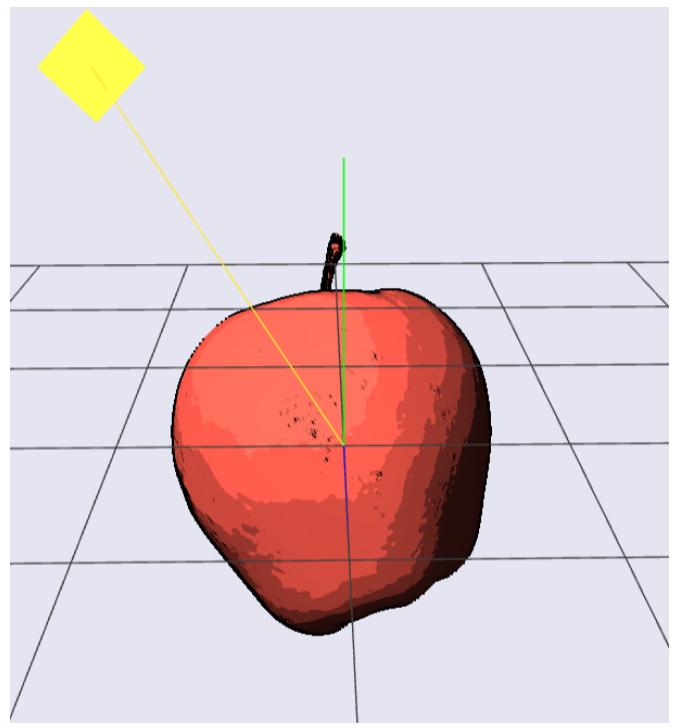
- 1) Multiple NPR Styles: Four distinct techniques implemented (exceeds requirement of 2)
- 2) Real-Time Performance: All effects run at ≥ 60 FPS on modern GPUs
- 3) Artistic Fidelity: Cross-hatching closely emulates Dürer's layered approach
- 4) Modularity: G-buffer pipeline enables easy addition of new post-processing effects

A. Comparison with Traditional Approaches

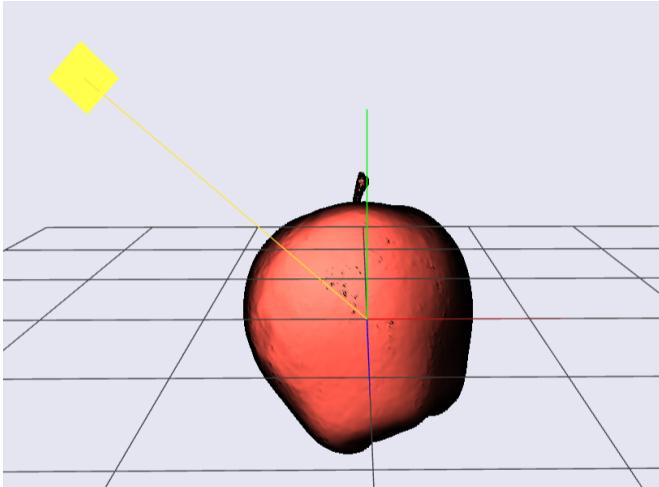
Our procedural cross-hatching differs from texture-based TAM (Tonal Art Maps) approaches [?]:



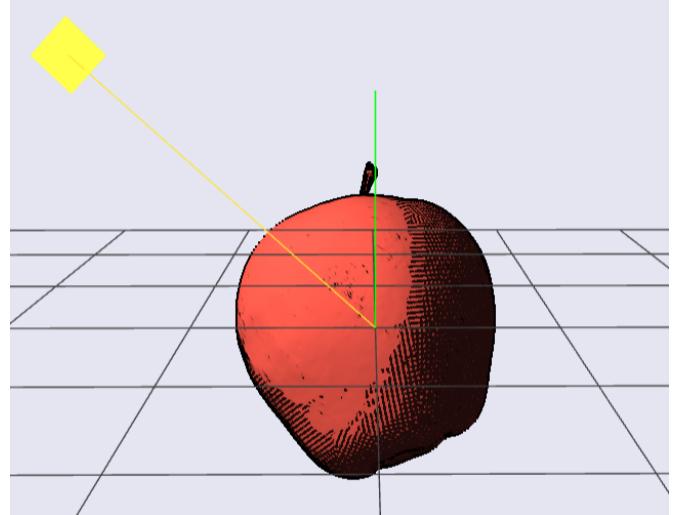
(a) Reference Phong



(b) Toon (4 bands)



(c) Edge Detection



(d) Cross-Hatching

Fig. 5: All NPR Techniques Applied to Apple Model (identical camera angle and lighting)

VII. Future Work

Potential enhancements:

- 1) Principal Curvature Alignment: Compute surface curvature directions to align hatching with form (requires geometry processing)
- 2) Temporal Coherence: Add temporal filtering to edge detection to reduce flickering during camera motion
- 3) Additional NPR Styles:
 - Watercolor simulation (wet-on-wet bleeding effects)

- Stippling (dot-based shading)
- Sketch lines (irregular, hand-drawn appearance)

- 4) Multi-Light Support: Extend to multiple light sources (currently single directional light)
- 5) Stroke Parameterization: Expose hatch angle, frequency, and threshold arrays to UI for fine-grained control

VIII. Conclusion

This project successfully implements multiple NPR techniques using modern GPU programming, demonstrat-

ing how classical artistic styles like Albrecht Dürer’s cross-hatching can be translated to real-time computer graphics. The modular architecture facilitates experimentation with different shading models, and the G-buffer approach enables sophisticated post-processing effects. All implementations achieve real-time performance while maintaining visual fidelity to their artistic inspirations.

The cross-hatching technique, in particular, shows how systematic analysis of traditional art techniques can inform algorithmic design. By decomposing Dürer’s layering strategy into discrete tone thresholds and directional patterns, we achieve a digital approximation that respects the historical medium while leveraging GPU parallelism.